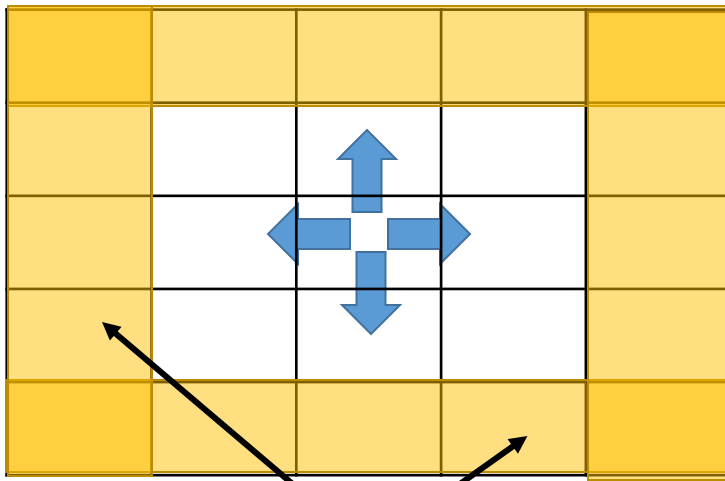


# Advanced MPI – RMA example

# Jacobi Solver

- The Jacobi method is an iterative algorithm for solving a system of linear equations.
- In the 2D model an approximation can be made by taking the average of the 4 neighbouring values (4 point stencil).



boundary

```
REAL A(0:n+1,0:n+1), B(1:n,1:n)
...
!Main Loop
DO WHILE(.NOT.converged)
  ! perform 4 point stencil
  DO j=1, n
    DO i=1, n
      B(i,j)=0.25*(A(i-1,j)+A(i+1,j)+A(i,j-1)+A(i,j+1))
    END DO
  END DO

  ! copy result back into array A
  DO j=1, n
    DO i=1, n
      A(i,j) = B(i,j)
    END DO
  END DO

  ...
  ! convergence test
END DO
```

# Jacobi in Parallel

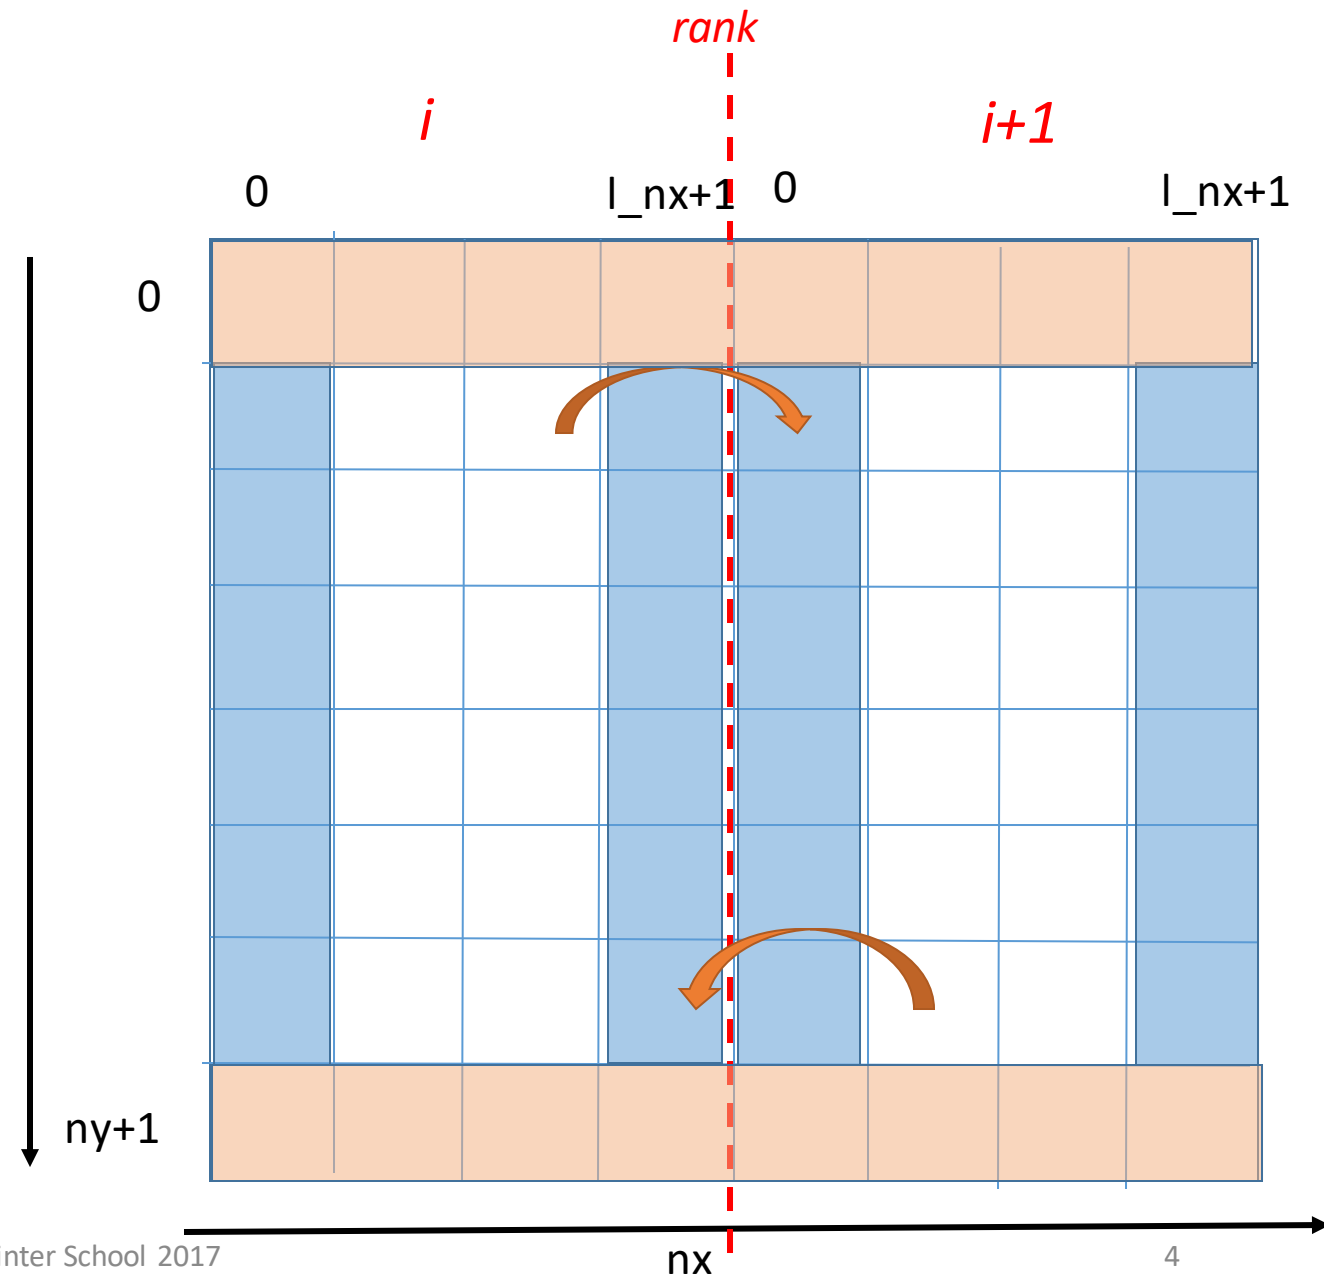
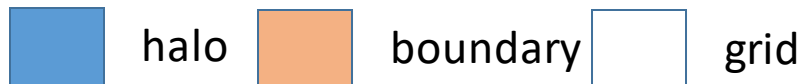
- For simplicity, can use 1D domain decomposition, dividing rows or columns of the grid among the MPI ranks.
- Since each rank will need data from neighbouring domains need to set up halo regions.
- For efficiency, exchange columns with Fortran and rows with C.
- As a first version, can use MPI\_Send and MPI\_Recv to exchange the data.

# Halo exchange -Fortran Version

We start with a data grid( $ny, nx$ ) which becomes  $grid(0:ny+1, 0:nx+1)$  when we include the boundaries.

Each local domain is  $(1..ny, 1..l_{nx})$  but with a halo region + boundaries we have  $(0..ny+1, 0..l_{nx}+1)$ .

Ranks 0 and size-1 need only 1 halo region, since they must store the left and right border conditions.

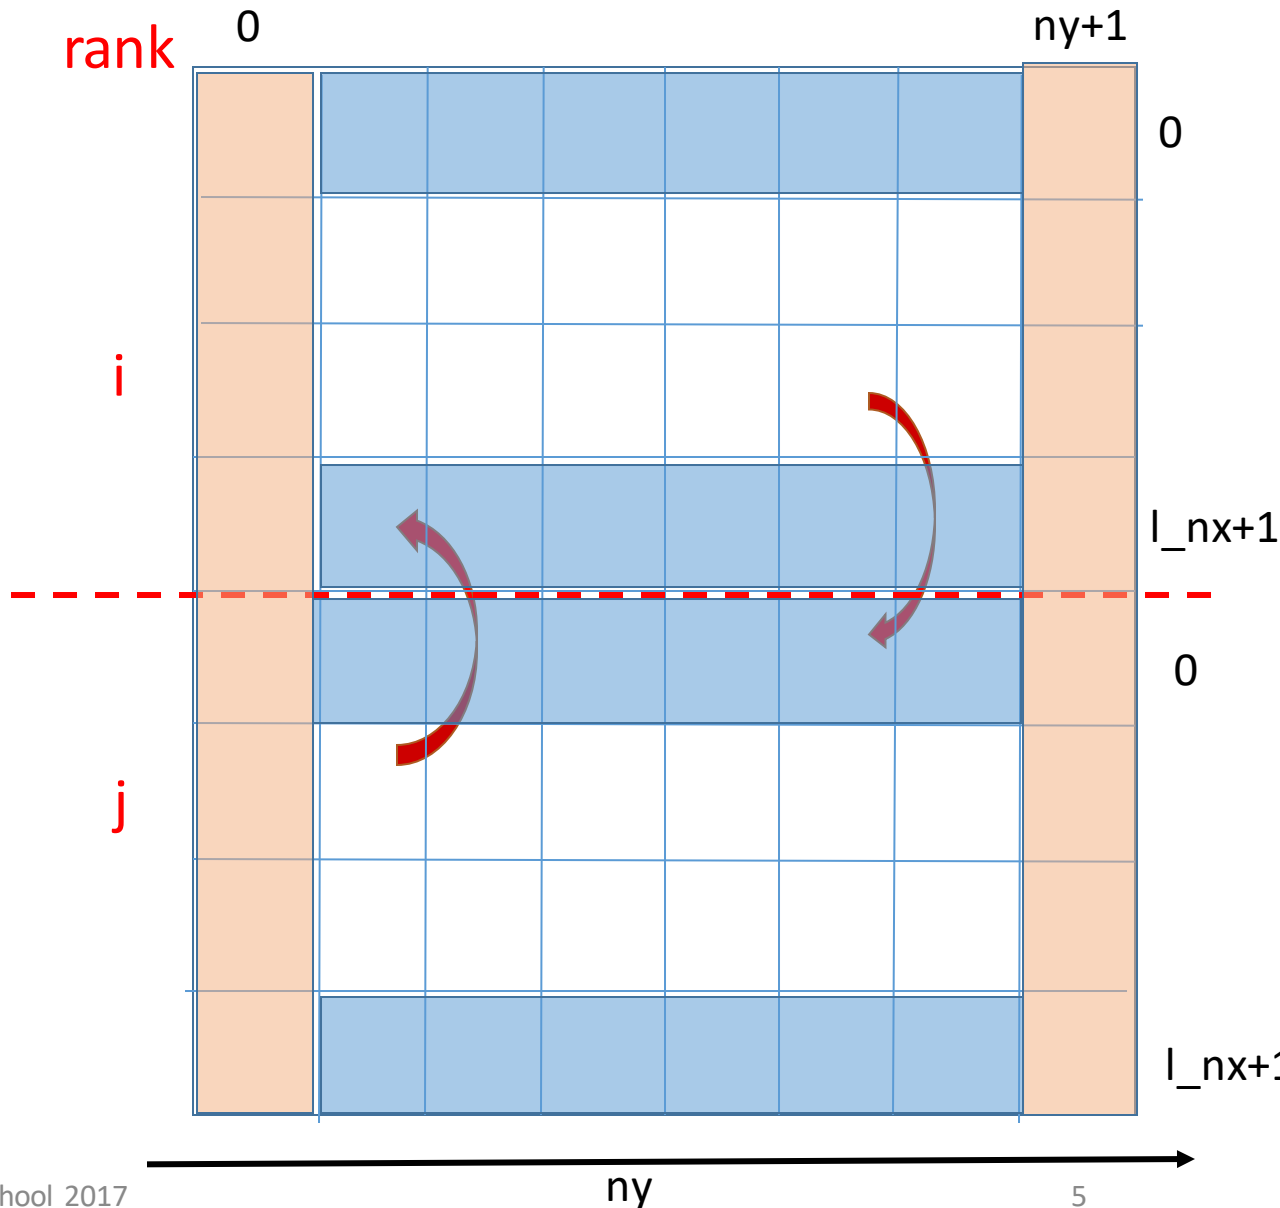
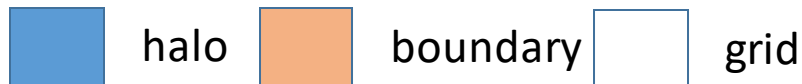


# Halo exchange C Version

For C we exchange rows instead of columns, such that the rows are divided among the MPI ranks.

MPI ranks 0 and size-1 need only 1 halo region (and only 2 transfers) because they contain the top and bottom boundary conditions.

Note we have inverted the nx and ny axes with respect to the FORTRAN version.



# Practical

1. Download the MPI C or Fortran source code:
  - git clone <https://gitlab.hpc.cineca.it/training/advanced-school.git>
2. Try first the serial code, for various grid sizes.
3. Look at the MPI version, compile, run and see how long it takes for various numbers of tasks:  
`mpicc -o jacobi jacobi.c OR mpiifort -o jacobi jacobi.f90`  
`mpirun -n 4 ./jacobi 10 10`
4. Replace the blocking send/recv with non-blocking send/recv and re-run. (optional: `mpi_sendrecv`).
5. Replace the non-blocking send/recv with `mpi_rma`.
6. Instead of fence synchronization you could try:
  1. PSCW (Post Start Complete Wait)
  2. Use lock and unlock.

*Inspired by the example at [http://www.archer.ac.uk/training/course-material/2016/09/160929\\_AdvMPI\\_EPCC/index.php](http://www.archer.ac.uk/training/course-material/2016/09/160929_AdvMPI_EPCC/index.php)*

